

BALTHAZAR

DAPHNE

LUCIE

BENES

LA LOUE

SIMON

# Handwritten Recognition

Neural Networks and Deep Learning

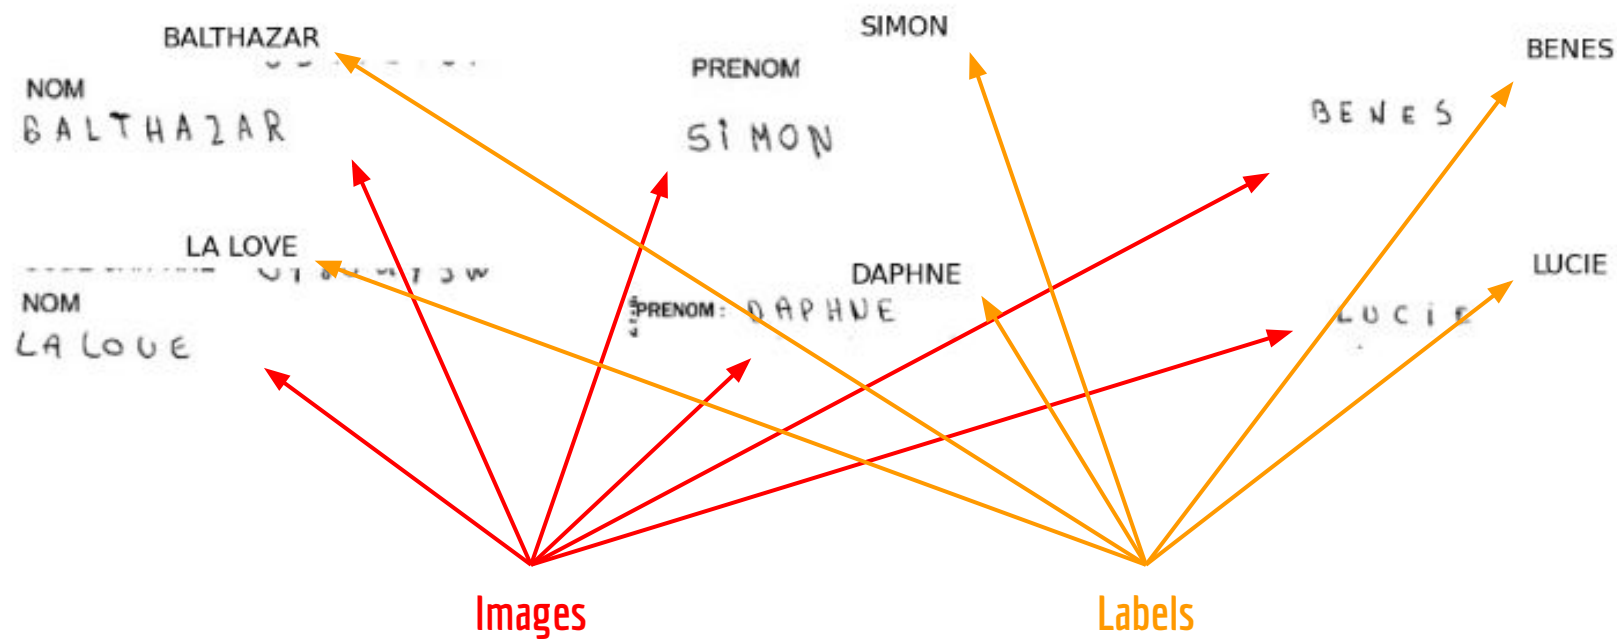
Andreu Gascón  
Mathias Lommel  
Pere Mayol



Context



# Context





The data



# The Data

**3 datasets** : Training / Validation / Test

For each dataset : **2 types of file**

- **.zip** file : containing the images

→ **330961** training / **41370** test and validation

- **.csv** file : linking each **image** with its **label**

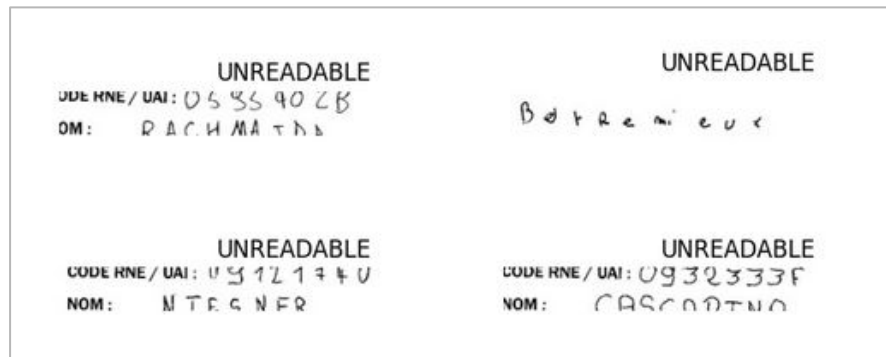


	TRAIN	FILENAME	IDENTITY
0	TRAIN_00001.jpg	BALTHAZAR	
1	TRAIN_00002.jpg	SIMON	
2	TRAIN_00003.jpg	BENES	
3	TRAIN_00004.jpg	LA LOVE	
4	TRAIN_00005.jpg	DAPHNE	

# The Data

For the labels :

- Delete existing **NaN** values
- Delete images labelled '**UNREADABLE**'



For the images :

- Cropping the image to a shape **256x64**
- Change type to **float**



# The Data

Finally...

- Code the labels into vectors
- Using a certain alphabet : " ABCDEFGHIJKLMNOPQRSTUVWXYZ-' "
- Padded into a length 24, with -1 values

WORD



[ 23, 16, 19, 5, -1, -1, -1, ..., -1 ]

Longest word found : 34 characters

Second longest : 24 characters

- We deleted the 34 character image from the training set


# Models



# Models

→ 1 model - 2 versions


→ Convolutional Recurrent Neural Network (**CRNN**)





# Models

→ 1 model - 2 versions

→ Convolutional Recurrent Neural Network (**CRNN**)



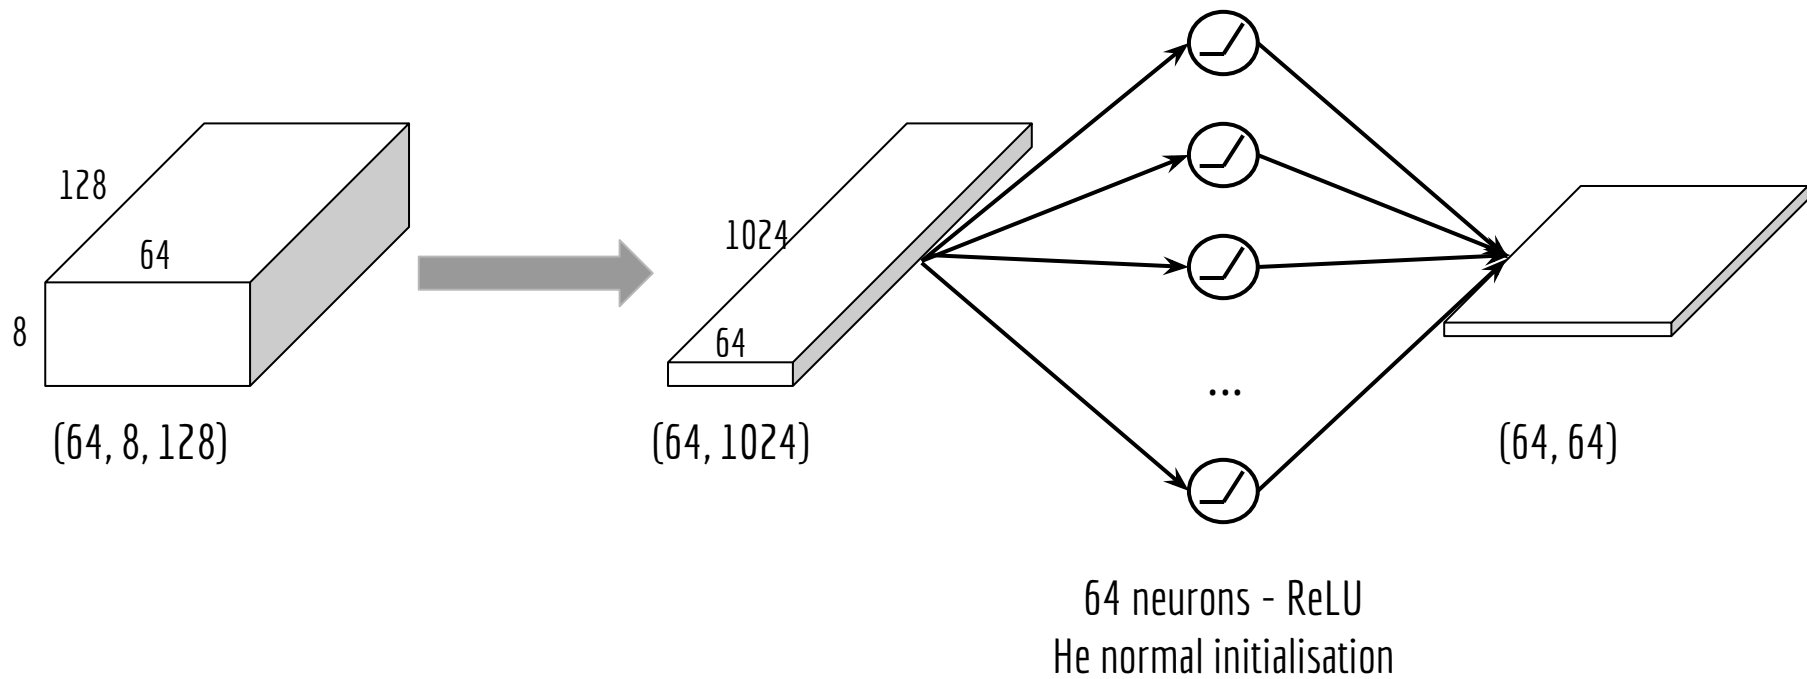
# Models



For each *convolutional* operation :

- *Batch normalisation*
- *He normal initialisation*


# Models



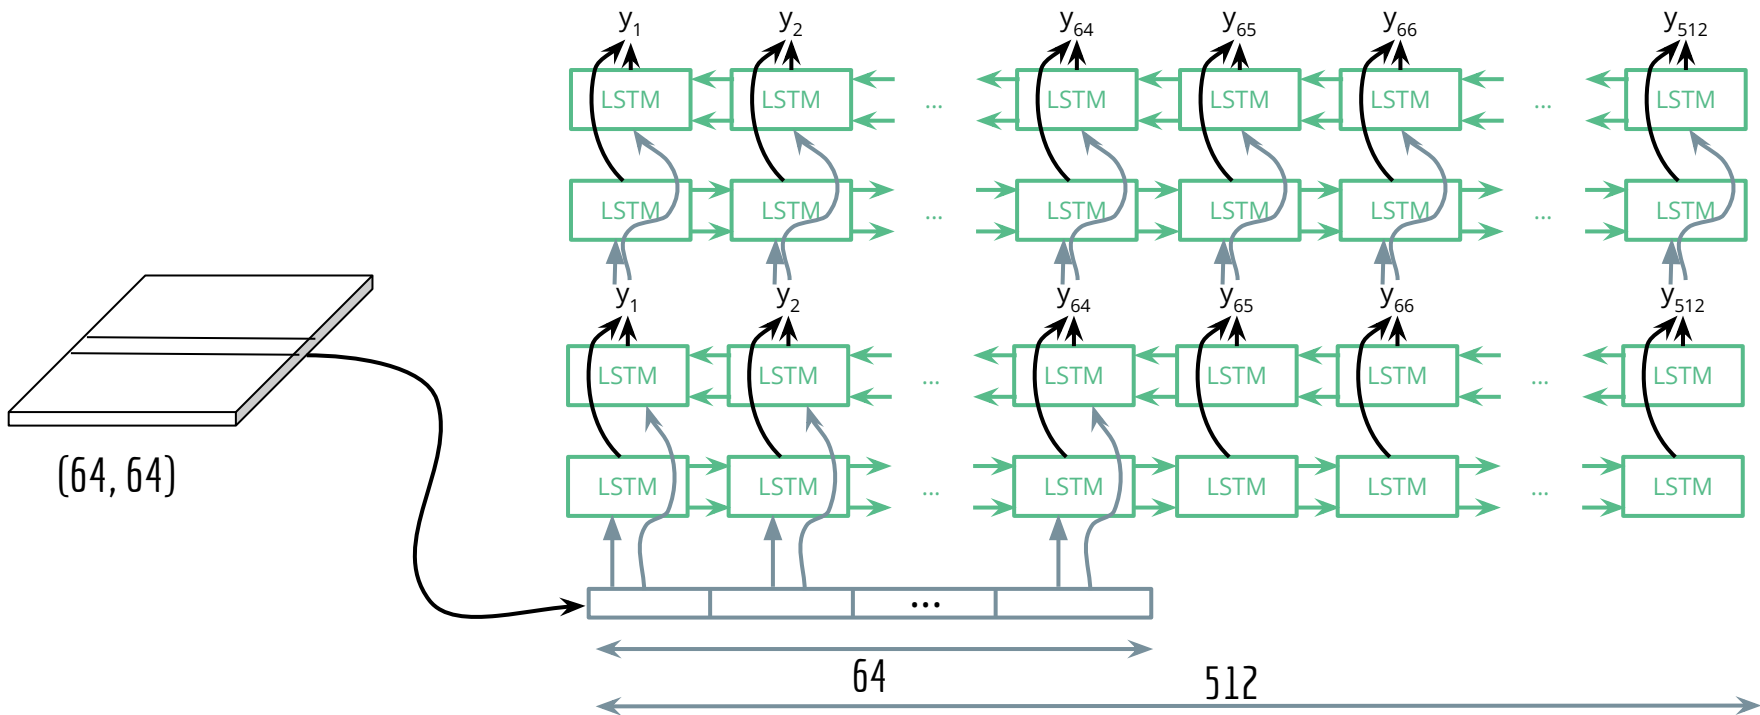

NOM  
BALHAZAR

CNN

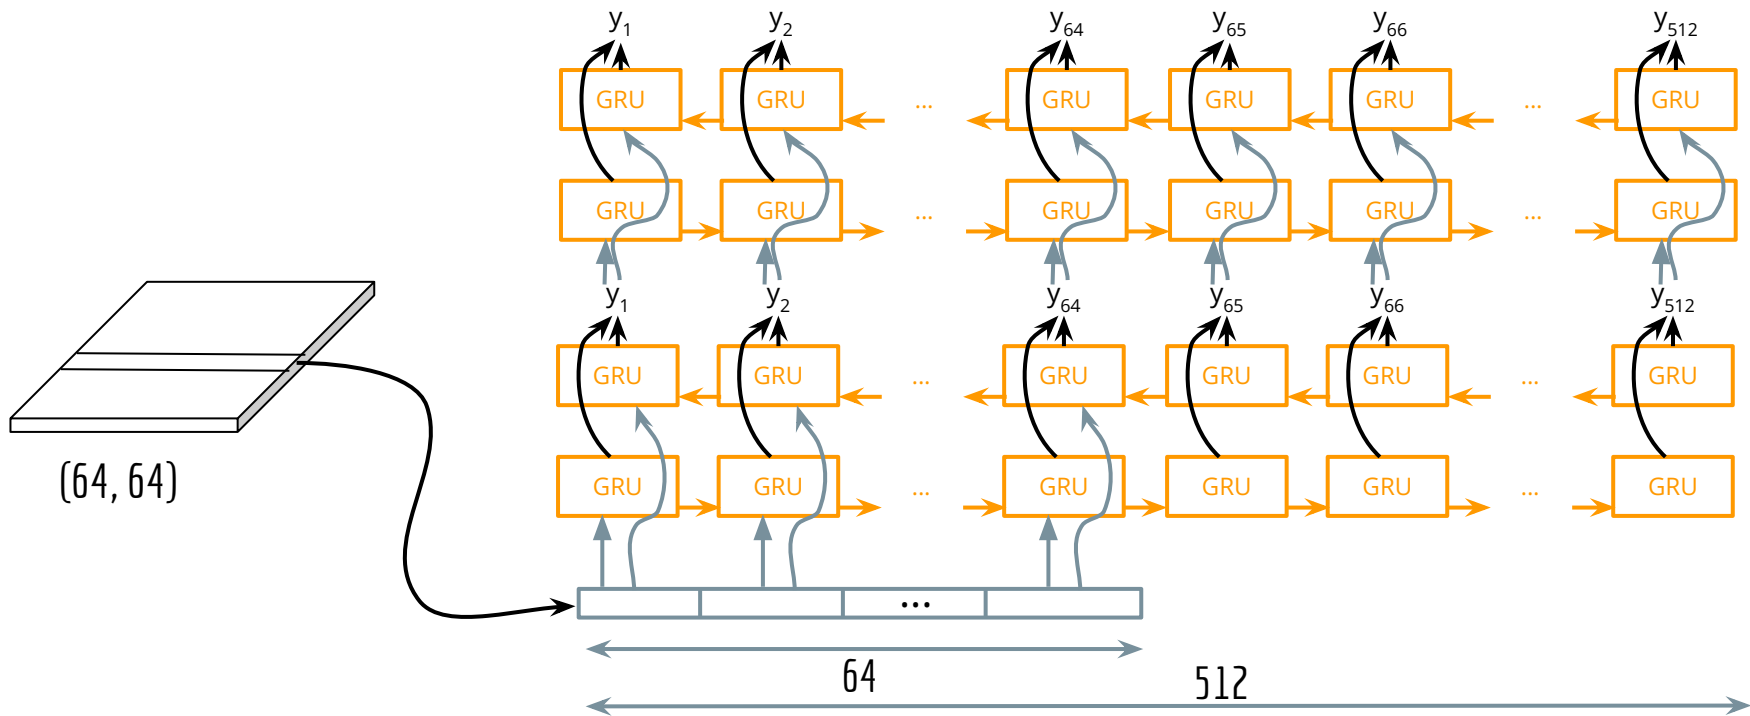

BALHAZAR





# Models



# Models




# Models




# Models

Why 64x30 ?



Probability distribution

alphabet = " ABCDEFGHIJKLMNOPQRSTUVWXYZ-"



$$Pr(\pi|\mathbf{x}) = \prod_t^T y_{\pi_t}^t \quad \text{Joint probability of a path} \quad \left\{ \begin{array}{l} \text{'WORD'} \\ \rightarrow \text{'WWWOORRRDDD'} \\ \rightarrow \text{'WOOOOORRRDDD'} \\ \rightarrow \dots \end{array} \right.$$

$$Pr(\mathbf{l}|\mathbf{x}) = \sum_{\pi \in \mathcal{B}^{-1}(\mathbf{l})} Pr(\pi|\mathbf{x}) \quad \text{Sum over all possible paths}$$


Connectionist Temporal  
Classification Loss

(CTC Loss)

$$\hookrightarrow \mathcal{L}_{CTC} = - \sum_{(\mathbf{x}, \mathbf{l}) \in \mathcal{Z}} \ln Pr(\mathbf{l}|\mathbf{x})$$



# Models



How to obtain the real label ?

By taking each max prob, we obtain :  $[23, 23, 16, 16, 16, 19, 19, 19, 19, 5, 5, -1, -1, -1, \dots, -1]$   
 $\longleftrightarrow$   
 64 characters

$[23, 23, -, 16, 16, 16, -, 19, 19, 19, 19, -, 5, 5, -, -1, -1, \dots, -1]$   $\longleftrightarrow$   $[23, 16, 19, 5, -1, -1, -1, \dots, -1]$

W	W	-	O	O	O	-	R	R	R	R	-	D	D
W	-		O			-		R			-	D	
W			O					R				D	



# Training

# Training setup

## *First training*

Size of the training sample : **64 000 images**

Size of the validation sample : **6 400 images**

Batch size : **128**

**Adam** optimizer

Learning rate of **0.001**

**10** epochs

## *Full training*

Size of the training sample : **300 800 images**

Size of the validation sample : **30 080 images**

Batch size : **128**

**Adam** optimizer

Learning rate of **0.001**

**10** epochs

# Tests



# Tests - first training

GRU model :

- Best word accuracy : **74.4%**
- Best letter accuracy : **89.9%**

6 689 822 parameters

No overfitting




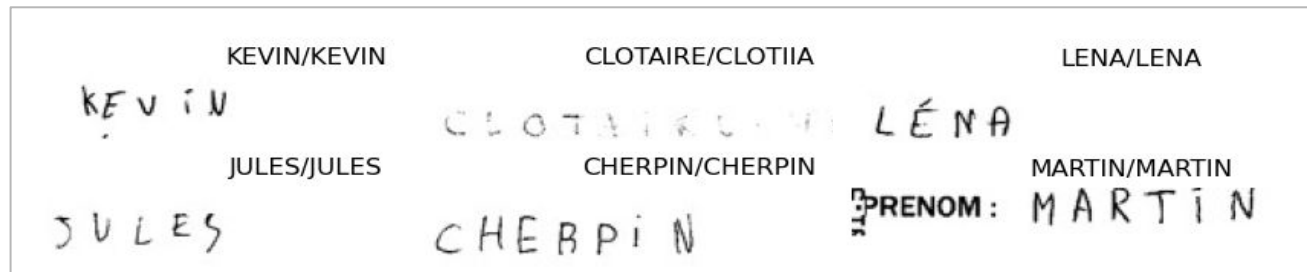
# Tests - first training

## LSTM model :

- Best word accuracy : **78.1%**
- Best letter accuracy : **91.3%**

8 856 606 parameters

An **overfitting** begins (more parameters)





# Tests - full training

## GRU model :

- Best word accuracy : **81 % (+6.6%)**
- Best letter accuracy : **92% (+2.1%)**

6 689 822 parameters

An **overfitting** begins to occur




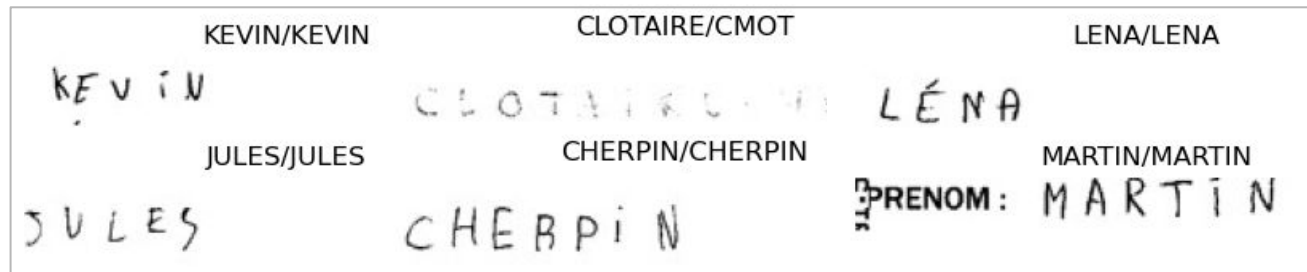
# Tests - full training

## LSTM model :

- Best word accuracy : **82.8% (+4.7%)**
- Best letter accuracy : **92.5% (+1.5%)**

8 856 606 parameters

Clear **overfitting**





# Tests - Our best model



Top **mispredicted** letters :

'N' predicted as 'M'

'H' predicted as 'M'

'O' predicted as 'U'

→ Similar letter shapes : the model has learnt something

On **mispredicted** words :

→ on average, **64.5%** of the letters are well predicted

Test on homemade images :

ANDREU

ANDREU

PIATHIAS

MATHIAS

PERE


P E R E

# Tests - Comparison of the models

	GRU (first)	LSTM (first)	GRU (full)	LSTM (full)
Word accuracy	74.4	78.1	81	82.8
Letter accuracy	89.9	91.3	92	92.5
Time of training	6'37"	7'21"	31'12"	33'31"
Overfitting ?	x	+	+	++
Number of parameters	6 689 822	8 856 606	6 689 822	8 856 606

# Tests - Our best model

*Kaggle Baseline Model*



*Correct characters predicted : 82.16%*  
*Correct words predicted: 69.10%*

# Problems

*Double letters problem*



*Changing the decoding function.*

*Poor performance for personal  
images*



*Preprocess the images properly.*

*Nan values for weights*



*Changing the dimension for  
`log_softmax(x,dim=2)`*

# Conclusion



Thanks for listening

